

Wikimedia image database project

Szymon Jaroszewicz

Overview

- Build a model which will tell if input image is a photo **taken outside**

Class 0



Class 1



- The hard part:
 - no training set available
 - need to prepare it yourself
- Idea:
 - work with unlabeled images
 - use publicly available models to generate class labels

- Wikimedia Foundation provided a dataset of 37 million images
 - scaled to 300px
 - captions in various languages
 - link to Wikipedia article
 - image embeddings
- https://huggingface.co/datasets/wikimedia/wit_base

- Several free models are available on [huggingface](#) (some require registration and token)
- You can:
 - Ask a Multimodal LLM the question about the image directly
 - Ask an ordinary LLM the question based on image caption
 - Use object detection model to find things in the image

To get started

- Collab spreadsheet which:
 - reads images from the Wikimedia dataset
 - asks the Gemma 3 model about the image

- https://home.ipipan.waw.pl/sj/AdvML/Wikimedia_project/test_set.tgz
- 500 images for your model to label

What to submit

- 1 A file with extension `.txt` with the format:

- first line: `nr indeksu` or another identifier
- following lines: predicted class labels for subsequent images
 - class 1 - outside
 - class 0 - not outside

- 2 Python code

Example file contents:

```
nr_indeksu
1
1
0
0
1
...
```